

Least Path Grade Transitions

Polypropylene has many potential grades each with unique characteristics. Each grade is created by a certain reaction environment. Changing this environment takes time and causes off-spec or wide-spec product to be produced.

To minimize this transition, it is best to move to similar reactions instead of very different reactions. If we look at a simple example, a planner would want to go from low melt flow grades to high ones and then back down in a sin wave. However, there are more complex attributes like if a grade is random of homo-polymer etc. To solve this we write a simple linear program. This is a real life application of the classic **Traveling Salesperson Problem** (TSP).

I would like to thank Evan Fields and [this](#) blog post for getting me started.

Loading Libraries

```
Plots.GRBackend()
```

```
• begin
•   using MultivariateStats
•   using JuMP
•   using Cbc
•   using Distances
•   using Distributions
•   using Random
•   using Plots
•   gr()
• end
```

Generating Data

First we need to set the transition costs. In the real world, we multi dimensional output space to calculate the transition matrix. In order to visualize this, we are going to use a 2D vector space to more easily visualize.

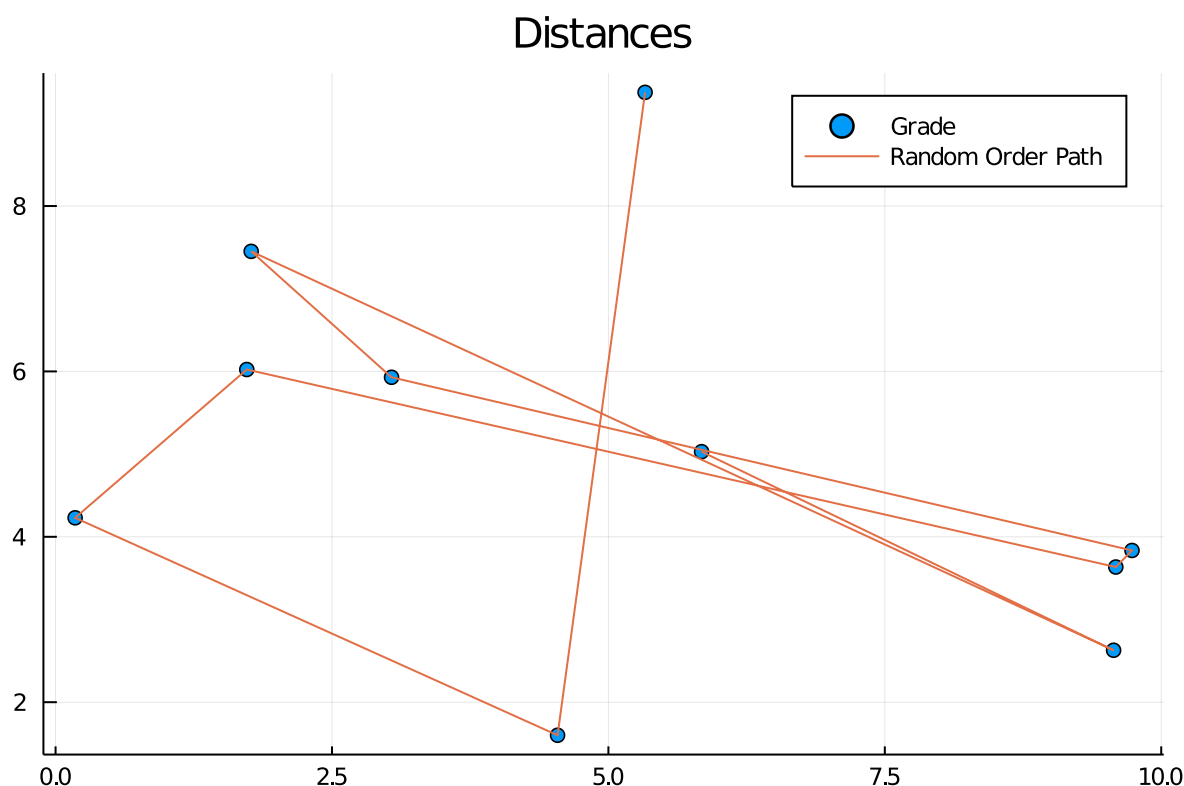
This difference is a bit academic as multi factored reasons for costs are projected down onto a matrix.

Float64[9.37466, 1.60006, 4.22956, 6.02298, 3.63458, 3.83491, 5.92912, 7.45181, 2.6:

```

• begin
•     len = 10
•     Random.seed!(42)
•     x_low_high = rand(Uniform{0.0, 10.0}, len)
•     y_low_high = rand(Uniform{0.0, 10.0}, len)
•
•
• end

```



```

• begin
•     plot(x_low_high, y_low_high, seriestype = :scatter, title = "Distances",
•         lab = "Grade")
•     plot!(x_low_high, y_low_high, lab = "Random Order Path")
• end

```

This is obviously not the best path through the production wheel.

Randomizing It

It's also not the only problem we face. In polypropylene reactions, the transition from high to low is different from the transition from low to high even for the same two grades. To simulate this we add some random noise.

Float64[13.7689, 1.72807, 5.12402, 7.83679, 4.29509, 4.57024, 7.68684, 10.2283, 2.9:

```


• begin
•     Random.seed!(42)

```

```

•
•
•   N = length(x_low_high)
•   x_perturbation = rand(Uniform(1.0, 1.5), N)
•   y_perturbation = rand(Uniform(1.0, 1.5), N)
•
•   x_high_low = x_low_high .* x_perturbation
•   y_high_low = y_low_high .* y_perturbation
• end

```



```

• begin
•   plot(x_low_high, y_low_high, seriestype = :scatter, title = "'Distances'", lab =
•     "Up Costs")
•   plot!(x_high_low, y_high_low, seriestype = :scatter, lab = "Down Costs")
• end

```

These perturbations exhibit the increased time of transitions when going in a different direction.

Product Wheel

Our product wheel should go from low to high and then back down. If our costs were the same on the way up and down, we could just use the same optimization but we are not guaranteed that it will be optimal.

Solving asymmetric problems like this is difficult so I am just going to break them up into two. If you had a traditional cost matrix that was asymmetric, you need to split it into two symmetric ones, one with the up costs on both sides of the diagonal and one with down costs. Since ours are artificial, I'm just

going to keep them in their separate forms. Just remember that most of the time we would need to split them.

So while we technically have a 'asymetric' problem with different values on the lower and upper triangles of the cost matrix, we split them into two and solve each individually.

```

• function cost_generator(d1, d2)
•     N = length(d1)
•     cost = zeros(N, N)
•     points = [[d1[i], d2[i]] for i ∈ 1:N]
•     for i ∈ 1:N, j ∈ 1:N
•         @inbounds cost[i, j] = euclidean(points[i], points[j])
•     end
•     return cost
• end;

```

```


cost_up =
10×10 Array{Float64,2}:
0.0      7.81479  7.28325  4.92055  7.14662  ...  4.13887  4.04852  7.96688  4.37509
7.81479  0.0      5.09448  5.24058  5.44347  ...  4.58209  6.47477  5.13287  3.66849
7.28325  5.09448  0.0      2.37202  9.43118  ...  3.32844  3.59417  9.52785  5.72217
4.92055  5.24058  2.37202  0.0      8.2148   ...  1.31273  1.42939  8.54331  4.23178
7.14662  5.44347  9.43118  8.2148   0.0      ...  6.9408  8.70209  1.0067  3.99769
7.0769   5.65565  9.56693  8.29994  0.248122 ...  7.01677  8.74919  1.21826  4.07199
4.13887  4.58209  3.32844  1.31273  6.9408   ...  0.0      1.98255  7.31736  2.94491
4.04852  6.47477  3.59417  1.42939  8.70209  ...  1.98255  0.0      9.17112  4.73951
7.96688  5.13287  9.52785  8.54331  1.0067   ...  7.31736  9.17112  0.0      4.43309
4.37509  3.66849  5.72217  4.23178  3.99769  ...  2.94491  4.73951  4.43309  0.0

```

```

• cost_up = cost_generator(x_low_high, y_low_high)

```




```

• heatmap(cost_up[end:-1:1, :]) # make it reverse b/c plots is annoying

```

```
cost_down =
10×10 Array{Float64,2}:
 0.0      12.0987  10.861    7.67785 ...  6.89728  5.98684  13.085    7.51687
12.0987   0.0      6.37278  7.13781 ...  6.30828  9.24894  8.66654  4.97657
10.861    6.37278   0.0      3.20165 ...  4.19563  5.395    14.1337  7.46367
 7.67785  7.13781   3.20165   0.0      ...  1.62844  2.39196  13.1975  5.87695
12.0421   8.99023  14.033   12.8075 ...  11.2119  13.6215  1.32225  6.93174
12.0099   9.34643  14.3071  13.0126 ...  11.4084  13.7658  1.62999  7.13642
 6.89728  6.30828   4.19563   1.62844 ...  0.0      2.98982  11.6439  4.28214
 5.98684  9.24894   5.395    2.39196 ...  2.98982  0.0      14.2131  6.86351
13.085    8.66654  14.1337  13.1975 ...  11.6439  14.2131  0.0      7.38644
 7.51687  4.97657   7.46367   5.87695 ...  4.28214  6.86351  7.38644  0.0
```

```
• cost_down = cost_generator(x_high_low, y_high_low)
```



```
• heatmap(cost_down[end:-1:1, :]) # make it reverse b/c plots is annoying
```

Classic Traveling Salesperson (TSP) problems are designed work for round trips. While we kind of want to a round trip with different costs associated with the up and down portions of the product wheel, we don't want to make our constraints too complex.

As a work around I just calculate one single up trip and one down trip. We still need to change the trips to be one way. To do this we add a trip that has astronomically high costs except for the start and end nodes which are zero. This way the end nodes are always 'next to each other' because they have the cheapest transition. The downside of this is that we have to pick the start and end nodes. To do this I picked the ones with the greatest cost (the maximum value of the cost matrix) to be the start and end.

We can then ignore this final transition when it comes to our analysis and the trip will look like it is one way.

```

• """
•     make_matrix_one_way(A::AbstractMatrix)
• Takes a Matrix, calculates the max possible transition cost and pads it with a vector
• of that cost. The highest transition cost values (farthest cities) are set to zero to
• allow for the MIP to complete a tour while maintaining uni-directional logic.
• """
• function make_matrix_one_way(A::AbstractMatrix; idx = nothing)
•     max_possible = sum(A)
•     if idx == nothing
•         idx = argmax(A)
•     end
•     id1, id2 = idx[1], idx[2]
•
•     shortcut_vec = fill(max_possible, size(A, 1) + 1)
•     shortcut_vec[[id1, id2, end]] .= 0.0
•     cost = hcat(A, shortcut_vec[1:end-1])
•     cost = vcat(cost, shortcut_vec')
•
•     return cost, idx
• end;
•

```

```

(11×11 Array{Float64,2}:
 0.0      7.81479  7.28325  4.92055 ...  7.96688  4.37509  488.357
 7.81479  0.0      5.09448  5.24058 ...  5.13287  3.66849  488.357
 7.28325  5.09448  0.0      2.37202    9.52785  5.72217  0.0
 4.92055  5.24058  2.37202  0.0      8.54331  4.23178  488.357
 7.14662  5.44347  9.43118  8.2148    1.0067   3.99769  488.357
 7.0769   5.65565  9.56693  8.29994 ...  1.21826  4.07199  0.0
 4.13887  4.58209  3.32844  1.31273    7.31736  2.94491  488.357
 4.04852  6.47477  3.59417  1.42939    9.17112  4.73951  488.357
 7.96688  5.13287  9.52785  8.54331    0.0      4.43309  488.357
 4.37509  3.66849  5.72217  4.23178    4.43309  0.0      488.357
 488.357  488.357  0.0      488.357 ...  488.357  488.357  0.0
, CartesianI

```

```

• cost1, _ = make_matrix_one_way(cost_up)

```

This allows us to simulate a one way trip if we remove the final index.

TSP Solution

I am using the **Dantzig-Fulkerson-Johnson formulation** to solve my problem.

It is easiest understood as thining of it in two parts. First, I wrote the `dfj` function to initialize the model `m` the Transition variable, and the objective function.

Secondly, the `dfj` function adds constraints to make sure that there are no zero cost joins back to themselves and then to make sure that each row and column have an exact value of 1. This makes it so each transition will happen once and only once.

While these constraints seem complete, there is actually one additional step that, most of the time, needs to be taken.

```

• "
•     dfj(A::AbstractArray, optimizer)
• Takes cost matrix and computes the first step of the optimization.
• "
• function dfj(A::AbstractArray, optimizer)
•     N, M = size(A)
•
•     m = Model(optimizer)
•     @assert N == M "Matrix must be square"
•     @variable(m, Transition[1:N, 1:M], Bin)
•
•     @objective(m, Min, sum(Transition .* A))
•
•     trans_iterator = 1:N
•     for i ∈ trans_iterator
•         @constraint(m, Transition[i, i] == 0)
•         # ^ No self linking
•         @constraint(m, sum(Transition[i, :]) == 1)
•         # ^ All must be used
•     end
•
•     for j ∈ trans_iterator
•         @constraint(m, sum(Transition[:, j]) == 1)
•     end
•
•     for i ∈ trans_iterator, j ∈ trans_iterator
•         @constraint(m, Transition[i, j] + Transition[j, i] ≤ 1)
•         # ^ makes it impossible to have to "same" tranistions
•         # on the diagonal or thought of another way, cities self joining with
•         # eachother
•     end
•
•     optimize!(m)
•     return m
• end;

```

Unfortunately, the constraints in the `dfj` function cause it to often return sub tours instead of a tour that hits every path. Think:

A -> B -> C -> A and D -> E -> D

This satisfies mini TSP problems but not the global one. We could exhaustively limit these issues but MIPs get very difficult to solve if we have too many constraints.

A way around this is to **Delayed Column Generation** We run the model, check if it has tours that are too small, limit those tours with a new constraint, and then run it again. We repeat this until we get a satisfactory result.

This provides the optimal solution with minimal constraints.

```

• "     delayed_column_gen(m)
• Takes a JuMP model (generated from the dfj()) and adds constraints as necessary to add
• complete tours.

```

```

• "
• function delayed_column_gen(m)
•     var = m.obj_dict[:Transition]
•     vals = value.(var)
•
•     N = size(vals, 1)
•     Q = [1]
•     while true
•         idx = argmax(vals[Q[end], :])
•         if idx == Q[1] # checks if tour is completed and short
•             break
•         else
•             push!(Q, idx)
•         end
•     end
•
•     if length(Q) < N
•         @show Q
•         @constraint(m, sum(var[Q, Q]) ≤ length(Q) - 1) # 'Bans' the current solution
•         optimize!(m)
•         delayed_column_gen(m)
•     else
•         return Q
•     end
•
• end;

```

```

• m = dfj(cost1, Cbc.Optimizer);

```

```

Int64[6, 5, 9, 2, 10, 1, 8, 7, 4, 3]


```

```

• begin
•     best_order_with_skip = delayed_column_gen(m)
•     skip_idx = argmax(best_order_with_skip)
•     best_order =
•         vcat(best_order_with_skip[1+skip_idx:end], best_order_with_skip[1:skip_idx-1])
• end

```


Distances



```

• begin
•     plot(
•         x_low_high,
•         y_low_high,
•         seriestype = :scatter,
•         title = "Distances",
•         series_annotations = [string(i) for i ∈ 1:10],
•         lab = "Grades"
•     )
•     plot!(x_low_high[best_order], y_low_high[best_order], lab= "Path")
• end

```

Next, lets put these into convenient functions to find the up and down tours.

```

• begin
•     function find_path(costs, optimizer)
•         costs, idx = make_matrix_one_way(costs)
•         m = dfj(costs, optimizer)
•
•         best_order = delayed_column_gen(m)
•         skip_idx = argmax(best_order)
•         if skip_idx == length(best_order)
•             best_order = best_order[1:end-1]
•         elseif skip_idx == 1
•             best_order = best_order[2:end]
•         else
•             best_order = vcat(best_order_with_skip[1+skip_idx:end],
•                               best_order_with_skip[1:skip_idx-1])
•         end
•         return best_order, idx
•     end
•     function find_path(costs, optimizer, start_end)
•
•         costs, idx = make_matrix_one_way(costs, idx=start_end)
•         m = dfj(costs, optimizer)

```

```

•
•
•     best_order = delayed_column_gen(m)
•     # return best_order
•     skip_idx = argmax(best_order)
•     if skip_idx == length(best_order)
•         best_order = best_order[1:end-1]
•     elseif skip_idx == 1
•         best_order = best_order[2:end]
•     else
•         best_order = vcat(best_order_with_skip[1+skip_idx:end],
•                           best_order_with_skip[1:skip_idx-1])
•     end
•     return best_order
• end
• end;

```

```
(Int64[6, 5, 9, 2, 10, 1, 8, 7, 4, 3], CartesianIndex(6, 3))
```

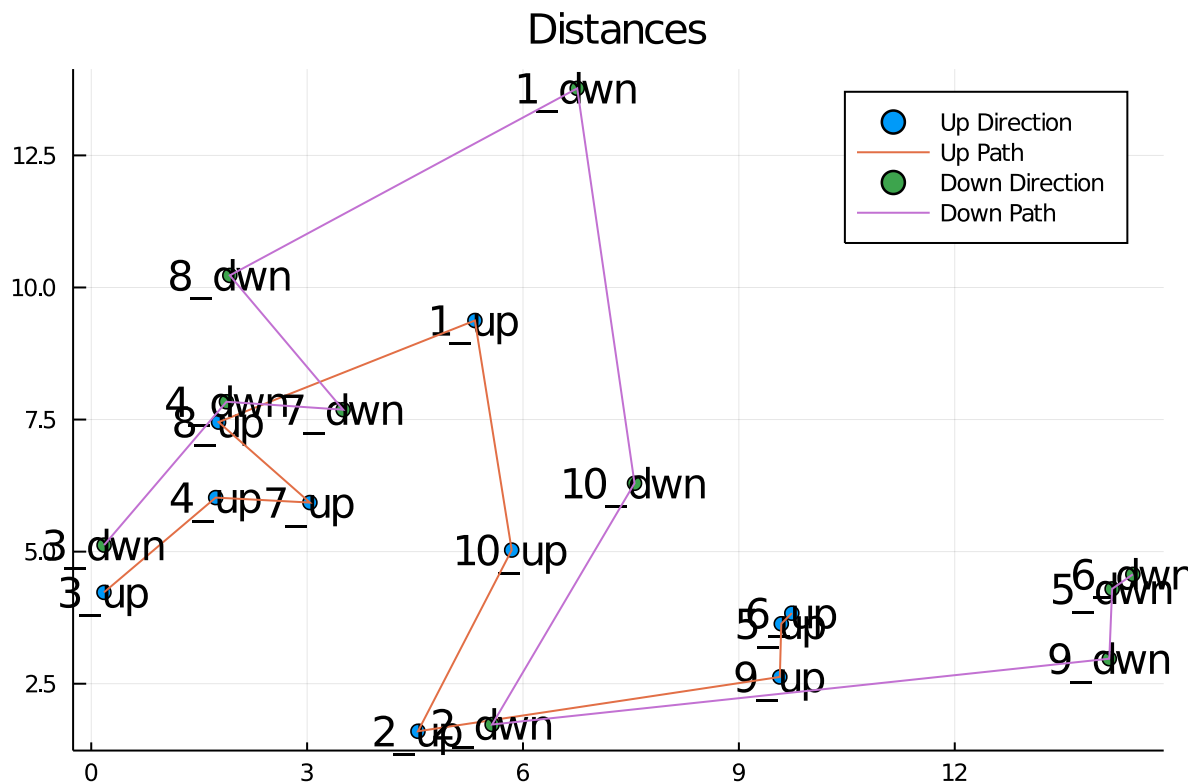
```
• up_order, idx = find_path(cost_up, Cbc.Optimizer)
```

Note in the bellow function call we are using `idx` which substitutes the maximum cost of the down matrix with the start and end nodes of the previous up path.

If instead we wanted to arbitrarily set the start and end points we could supply them directly for both the up and down paths.

```
down_order = Int64[6, 5, 9, 2, 10, 1, 8, 7, 4, 3]
```

```
• down_order = find_path(cost_down, Cbc.Optimizer, idx)
```



```

• begin
•     plot(
•         x_low_high,

```

```

•      y_low_high,
•      seriestype = :scatter,
•      title = "Distances",
•      series_annotations = [string(i) * "_up" for i ∈ 1:10],
•      lab = "Up Direction"
•    )
•    plot!(x_low_high[up_order], y_low_high[up_order], lab = "Up Path")
•
•    plot!(
•      x_high_low,
•      y_high_low,
•      seriestype = :scatter,
•      title = "Distances",
•      series_annotations = [string(i) * "_down" for i ∈ 1:10],
•      lab = "Down Direction"
•    )
•    plot!(x_high_low[down_order], y_high_low[down_order], lab = "Down Path")
•
• end

```

```

(
1:  Int64[6, 5, 9, 2, 10, 1, 8, 7, 4, 3]
2:  Int64[6, 5, 9, 2, 10, 1, 8, 7, 4, 3]
)

```

```
• (up_order, down_order)
```

With a seed of 42, we get the same values for up and down. This will often be the case but they can potentially be different.

Final Steps

Often we will just have a cost matrix with no corresponding x, y vectors that we can plot the path on. (Remember above we calculated the costs with the Euclidean function.)

We use a technique called **Classical Multidimensional Scaling (MDS)** to project the distance matrix into two vectors.

The orientation and scaling are different but relative positions are maintained.

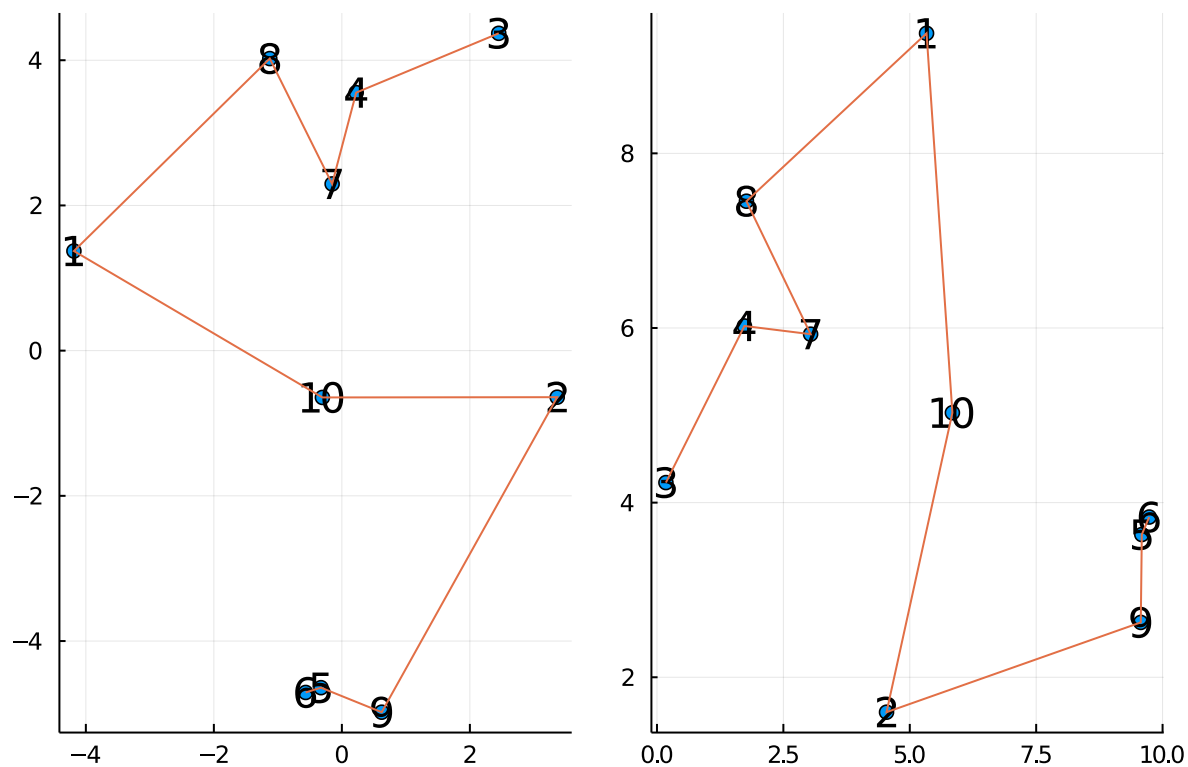
This captures the non-linear transform of the Euclidean formula very nicely. In real world grade transitions it may not perform exactly as well but it should at least maintain rough differences.

```

est =
2×10 Array{Float64,2}:
 1.37212  -0.64112  4.37075  3.55408  -4.64228  ...  4.02294  -4.97994  -0.644655
-4.18653   3.36448  2.45078  0.223783  -0.32633  ... -1.12652   0.622047  -0.304009

```

```
• est = classical_mds(cost_up, 2)
```



```

• begin
•   d1, d2 = (2, 1)
•   l = @layout [a b]
•   p1 = plot(est[d1, :], est[d2, :],
•             seriotype = :scatter,
•             series_annotations = [string(i) for i ∈ 1:10])
•   plot!(est[d1, :][up_order], est[d2, :][up_order])
•   p2 = plot(x_low_high, y_low_high,
•             seriotype = :scatter,
•             series_annotations = [string(i) for i ∈ 1:10])
•   plot!(x_low_high[up_order], y_low_high[up_order])
•
•   plot(p1, p2, layout = l, legend = false)
• end
•

```